

▼ End-to-end Multil-class Dog Breed Classification

This notebook builds an end-to-end multi-class image classifier using TensorFlow 2.0 and TensorFlow Hub.

1. Problem

Identifying the breed of a dog given an image of a dog.

2. Data

The data we're using is from Kaggle's dog breed identification competition.

<https://www.kaggle.com/c/dog-breed-identification/data>

3. Evaluation

The evaluation is a file with prediction probabilities for each dog breed of each test image.

<https://www.kaggle.com/c/dog-breed-identification/overview/evaluation>

4. Features

Some information about the data:

- We're dealing with images (unstructured data) so it's probably best we use deep learning/transfer learning.
- There are 120 breeds of dogs (this means there are 120 different classes).
- There are around 10,000+ images in the training set (these images have labels)
- There are around 10,000+ images in the test set (these images have no labels, because we'll want to predict them)

```
# Use the '-d' parameter as the destination for where the files should go
#!unzip "drive/MyDrive/dog Vision/dog-breed-identification.zip" -d "drive/MyDrive/dog Visi
```

▼ Get our workspace ready

- Import TenserFlow
- Import TenserFlow Hub
- we sure we use GPU

```
# Import TF 2.x
try:
    # %tensorflow_version only exists in Colab
```

```
%tensorflow_version 2.x
except Exception:
    pass

Colab only includes TensorFlow 2.x; %tensorflow_version has no effect.

# import necessary tools
import tensorflow as tf
import tensorflow_hub as hub
print("TF version:",tf.__version__)
print("TF Hub version:", hub.__version__)
# check for Gpu
print("GPU", "yes" if tf.config.list_physical_devices("GPU") else "not")

TF version: 2.11.0
TF Hub version: 0.12.0
GPU yes
```

▼ Getting our data ready (tuning into Tensors)

All machine learning models, our data has to be in numerical format so change it into numbers

Tuning our Image into Tensors (numerical representation)

```
from google.colab import drive
drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.remount()
```

```
#Lets start by accessing our data and cheching out the labels
import pandas as pd
labels_csv = pd.read_csv("drive/MyDrive/dog Vision/labels.csv")
print(labels_csv.describe())
print(labels_csv.head())
```

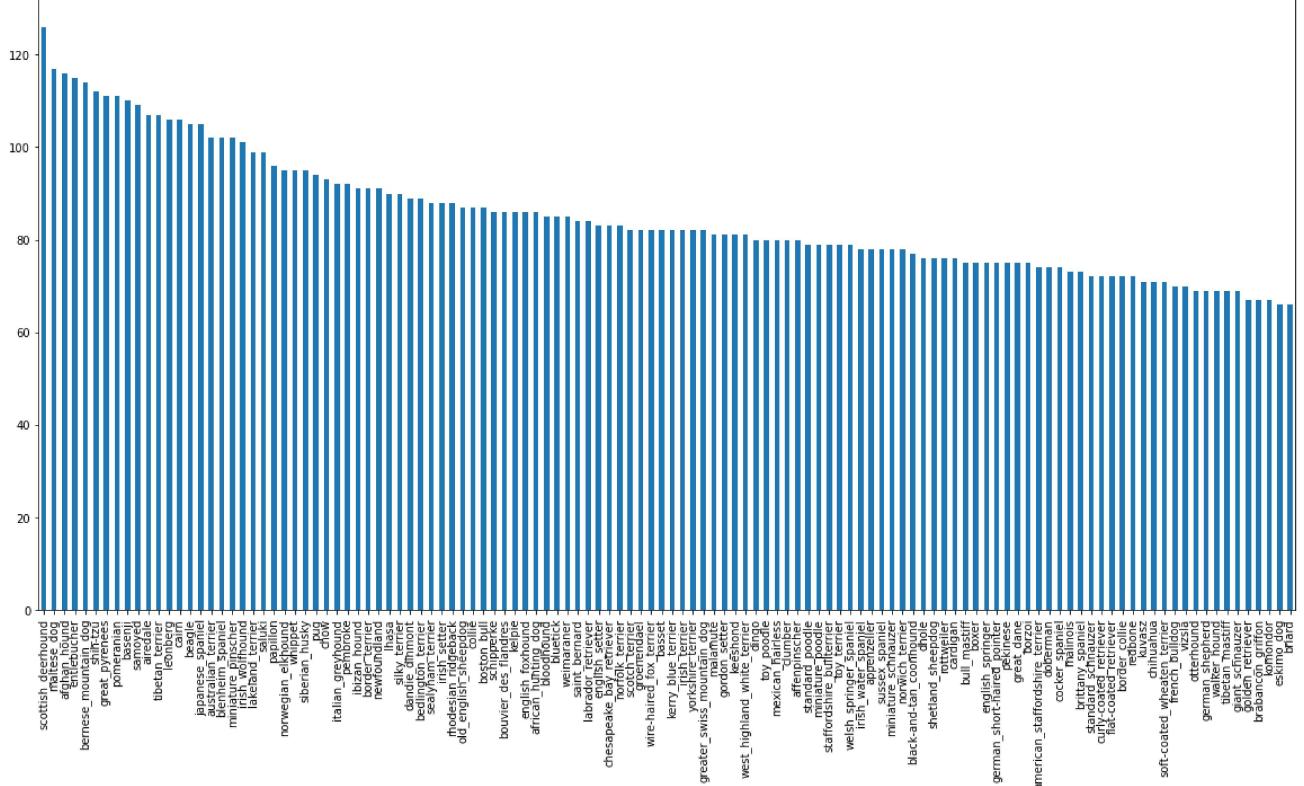
		id	breed
count		10222	10222
unique		10222	120
top	000bec180eb18c7604dcecc8fe0dba07	scottish_deerhound	
freq		1	126
		id	breed
0	000bec180eb18c7604dcecc8fe0dba07	boston_bull	
1	001513dfcb2ffafc82cccf4d8bbaba97	dingo	
2	001cdf01b096e06d78e9e5112d419397	pekinese	
3	00214f311d5d2247d5dfe4fe24b2303d	bluetick	
4	0021f9ceb3235effd7fcde7f7538ed62	golden_retriever	

```
labels_csv.head()
```

	id	breed
0	000bec180eb18c7604dcecc8fe0dba07	boston_bull
1	001513dfcb2ffafc82cccf4d8bbaba97	dingo
2	001cdf01b096e06d78e9e5112d419397	pekinese
3	00214f311d5d2247d5dfe4fe24b2303d	bluetick
4	0021f9ceb3235effd7fcde7f7538ed62	golden_retriever

How many images are there of each breed

```
labels_csv["breed"].value_counts().plot.bar(figsize= (20, 10));
```



```
from IPython.display import display, Image
```

```
labels_csv["breed"].value_counts().median()
```

```
82.0
```

```
# let's view an image
from IPython.display import Image
Image("drive/MyDrive/dog Vision/train/001cdf01b096e06d78e9e5112d419397.jpg")
```



▼ Getting images and their label

Let's get a list of all our image file pathnames

```
# Create path name from image ID's
filenames = ["drive/MyDrive/dog Vision/train/" + fname + ".jpg" for fname in labels_csv["i"]

# Check the first 10
filenames[:10]

['drive/MyDrive/dog Vision/train/000bec180eb18c7604dcecc8fe0dba07.jpg',
 'drive/MyDrive/dog Vision/train/001513dfcb2ffafc82cccf4d8bbaba97.jpg',
 'drive/MyDrive/dog Vision/train/001cdf01b096e06d78e9e5112d419397.jpg',
 'drive/MyDrive/dog Vision/train/00214f311d5d2247d5dfe4fe24b2303d.jpg',
 'drive/MyDrive/dog Vision/train/0021f9ceb3235effd7fcde7f7538ed62.jpg',
 'drive/MyDrive/dog Vision/train/002211c81b498ef88e1b40b9abf84e1d.jpg',
 'drive/MyDrive/dog Vision/train/00290d3e1fdd27226ba27a8ce248ce85.jpg',
 'drive/MyDrive/dog Vision/train/002a283a315af96ea0e28e7163b21b.jpg',
 'drive/MyDrive/dog Vision/train/003df8b8a8b05244b1d920bb6cf451f9.jpg',
 'drive/MyDrive/dog Vision/train/0042188c895a2f14ef64a918ed9c7b64.jpg']

# Check whether number of filenames matches number of actual image files
import os
```

```
if len(os.listdir("drive/MyDrive/dog Vision/train")) == len(filenames):
    print("Filenames match")
else:
    print("not match")

    Filenames match

# one more check
Image (filenames[9000])
```



```
labels_csv["breed"][9000]

'tibetan_mastiff'
```

since we've now get our training image filenames in a list let's prepare our labels

```
import numpy as np
labels = labels_csv["breed"].to_numpy() # convert labels column to NumPy array
labels[:10]
```

```
array(['boston_bull', 'dingo', 'pekinese', 'bluetick', 'golden_retriever',
       'bedlington_terrier', 'bedlington_terrier', 'borzoi', 'basenji',
       'scottish_deerhound'], dtype=object)

len(labels)
10222

# see if number of labels matches the number of filenames
if len(labels) == len(filenames):
    print("Number of labels matches number of filenames")
else:
    print("not match check data")

Number of labels matches number of filenames

# Find the unique label value
unique_breeds = np.unique(labels)
len(unique_breeds)

120

unique_breeds

array(['affenpinscher', 'afghan_hound', 'african_hunting_dog', 'airedale',
       'american_staffordshire_terrier', 'appenzeller',
       'australian_terrier', 'basenji', 'basset', 'beagle',
       'bedlington_terrier', 'bernese_mountain_dog',
       'black-and-tan_coonhound', 'blenheim_spaniel', 'bloodhound',
       'bluetick', 'border_collie', 'border_terrier', 'borzoi',
       'boston_bull', 'bouvier_des_flandres', 'boxer',
       'brabancon_griffon', 'briard', 'brittany_spaniel', 'bull_mastiff',
       'cairn', 'cardigan', 'chesapeake_bay_retriever', 'chihuahua',
       'chow', 'clumber', 'cocker_spaniel', 'collie',
       'curly-coated_retriever', 'dandie_dinmont', 'dhole', 'dingo',
       'doberman', 'english_foxhound', 'english_setter',
       'english_springer', 'entlebucher', 'eskimo_dog',
       'flat-coated_retriever', 'french_bulldog', 'german_shepherd',
       'german_short-haired_pointer', 'giant_schnauzer',
       'golden_retriever', 'gordon_setter', 'great_dane',
       'great_pyrenees', 'greater_swiss_mountain_dog', 'groenendael',
       'ibican_hound', 'irish_setter', 'irish_terrier',
       'irish_water_spaniel', 'irish_wolfhound', 'italian_greyhound',
       'japanese_spaniel', 'keeshond', 'kelpie', 'kerry_blue_terrier',
       'komondor', 'kuvasz', 'labrador_retriever', 'lakeland_terrier',
       'leonberg', 'lhasa', 'malamute', 'malinois', 'maltese_dog',
       'mexican_hairless', 'miniature_pinscher', 'miniature_poodle',
       'miniature_schnauzer', 'newfoundland', 'norfolk_terrier',
       'norwegian_elkhound', 'norwich_terrier', 'old_english_sheepdog',
       'otterhound', 'papillon', 'pekinese', 'pembroke', 'pomeranian',
       'pug', 'redbone', 'rhodesian_ridgeback', 'rottweiler',
       'saint_bernard', 'saluki', 'samoyed', 'schipperke',
       'scotch_terrier', 'scottish_deerhound', 'sealyham_terrier',
       'shetland_sheepdog', 'shih-tzu', 'siberian_husky', 'silky_terrier',
```



```
False, False, False, False, False, False, False, False,
False, False, False])

len (boolean_labels)

10222

# Turning boolean array into numbers

print(labels[0]) # original label
print(np.where(unique_breeds == labels[0])[0][0]) # index where label occurs
print(boolean_labels[0].argmax()) # index where label occurs in boolean array
print(boolean_labels[0].astype(int)) # there will be a 1 where the sample label occurs

boston_bull
19
19
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]

print(labels[2])
print(boolean_labels[2].astype(int))

pekingese
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
```

▼ Creating our own validation set

dataset from kaggle doesn't come with validation set we're going to create our own

```
# set x and y
X = filenames
y = boolean_labels

len(filenames)
```

10222

we're going to start experimenting with 1000 images and increase as needed

```
# set number of images to use for experimenting NUM_IMAGES:
NUM_IMAGES = 1000 #@param {type:"slider", min:1000, max:10000, step:1000}
```

1000

```

# Import train_test_split from Scikit-Learn
from sklearn.model_selection import train_test_split

# Split them into training and validation using NUM_IMAGES
X_train, X_val, y_train, y_val = train_test_split(X[:NUM_IMAGES],
                                                y[:NUM_IMAGES],
                                                test_size=0.2,
                                                random_state=42)

len(X_train), len(y_train), len(X_val), len(y_val)
(800, 800, 200, 200)

#let's have a geez at the training data
X_train[:5], y_train[:2]

([['drive/MyDrive/dog Vision/train/00bee065dcec471f26394855c5c2f3de.jpg',
  'drive/MyDrive/dog Vision/train/0d2f9e12a2611d911d91a339074c8154.jpg',
  'drive/MyDrive/dog Vision/train/1108e48ce3e2d7d7fb527ae6e40ab486.jpg',
  'drive/MyDrive/dog Vision/train/0dc3196b4213a2733d7f4bdcd41699d3.jpg',
  'drive/MyDrive/dog Vision/train/146fbfac6b5b1f0de83a5d0c1b473377.jpg'],
 [array([False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False]),
 array([False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, True, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False])
])

```

▼ turning images into tensors

to preprocess our images into tensors we're going to write a function we do:

1. Take a image file path as input

2. Use TensorFlow to read the file and save it to variable 'image'
3. Turn our image (a jpg) into Tensors
4. resize the image to be a shape of (224, 224)
5. return a modified image

```
# convert a image into numpy
from matplotlib.pyplot import imread
image = imread(filenames[42])
image.shape

(257, 350, 3)

# turn image into Tensor
tf.constant(image)[:2]

<tf.Tensor: shape=(2, 350, 3), dtype=uint8, numpy=
array([[[ 89, 137,  87],
       [ 76, 124,  74],
       [ 63, 111,  59],
       ...,
       [ 76, 134,  86],
       [ 76, 134,  86],
       [ 76, 134,  86]],

      [[ 72, 119,  73],
       [ 67, 114,  68],
       [ 63, 111,  63],
       ...,
       [ 75, 131,  84],
       [ 74, 132,  84],
       [ 74, 131,  86]]], dtype=uint8)>

image.max(), image.min()

(255, 0)
```

Now we've seen what image looks like as a Tensor let's make a function to preprocess them

1. Take a image file path as input
2. Use TensorFlow to read the file and save it to variable 'image'
3. Turn our image (a jpg) into Tensors
4. Normalize our image (convert color channels value from 0-255, to 0-1)
5. resize the image to be a shape of (224, 224)
6. return a modified image

```
# Define image size
IMG_SIZE = 224

# Create a function for preprocessing images
def process_image(image_path, img_size=IMG_SIZE):
```

```

"""
Takes an image file path and turns the image into a Tensor.
"""

# Read in an image file
image = tf.io.read_file(image_path)
# Turn the jpeg image into numerical Tensor with 3 colour channels (Red, Green, Blue)
image = tf.image.decode_jpeg(image, channels=3)
# Convert the colour channel values from 0-255 to 0-1 values
image = tf.image.convert_image_dtype(image, tf.float32)
# Resize the image to our desired value (224, 224)
image = tf.image.resize(image, size=[IMG_SIZE, IMG_SIZE])

return image

```

▼ Turning our data into batches

why turn our data into batches?

as we're trying to process 10k+ images they all might not fit in the memory

so that's why we do about 32(this is batch size) image at a time you can manually adjust the batch size if need

in order to use Tensorflow effectively we need our data in the form of Tensor type which look like this: '(image, label)

```

# Create a simple function to return a tuple (image, label)
def get_image_label(image_path, label):
    """
    Takes an image file path name and the associated label,
    processes the image and returns a tuple of (image, label).
    """

    image = process_image(image_path)
    return image, label

# demo of above code
process_image(X[42], tf.constant(y[42]))

<tf.Tensor: shape=(224, 224, 3), dtype=float32, numpy=
array([[ [0.3264178 , 0.5222886 , 0.3232816 ],
       [0.2537167 , 0.44366494, 0.24117757],
       [0.25699762, 0.4467087 , 0.23893751],
       ...,
       [0.29325107, 0.5189916 , 0.3215547 ],
       [0.29721776, 0.52466875, 0.33030328],
       [0.2948505 , 0.5223015 , 0.33406618]],

      [[0.25903144, 0.4537807 , 0.27294815],
       [0.24375686, 0.4407019 , 0.2554778 ],
       [0.2838985 , 0.47213382, 0.28298813],
       ...,
       [0.2785345 , 0.5027992 , 0.31004712],
       [0.28428748, 0.5108719 , 0.32523635],
```

```

[0.28821915, 0.5148036 , 0.32916805]],

[[0.20941195, 0.40692952, 0.25792548],
 [0.24045378, 0.43900946, 0.2868911 ],
 [0.29001117, 0.47937486, 0.32247734],
 ...,
 [0.26074055, 0.48414773, 0.30125174],
 [0.27101526, 0.49454468, 0.32096273],
 [0.27939945, 0.5029289 , 0.32934693]],

...,

[[0.00634795, 0.03442048, 0.0258106 ],
 [0.01408936, 0.04459917, 0.0301715 ],
 [0.01385712, 0.04856448, 0.02839671],
 ...,
 [0.4220516 , 0.39761978, 0.21622123],
 [0.47932503, 0.45370543, 0.2696505 ],
 [0.48181024, 0.45828083, 0.27004552]],

[[0.00222061, 0.02262166, 0.03176915],
 [0.01008397, 0.03669046, 0.02473482],
 [0.00608852, 0.03890046, 0.01207283],
 ...,
 [0.36070833, 0.33803678, 0.16216145],
 [0.42499566, 0.3976801 , 0.21701711],
 [0.4405433 , 0.4139589 , 0.23183356]],

[[0.05608025, 0.06760229, 0.10401428],
 [0.05441074, 0.07435255, 0.05428263],
 [0.04734282, 0.07581793, 0.02060942],
 ...,
 [0.3397559 , 0.31265694, 0.14725602],
 [0.387725 , 0.360274 , 0.18714729],
 [0.43941984, 0.41196886, 0.23884216]]], dtype=float32)>

```

Now we've got a way to turn our data into tuples of tensors in form (image, label), let's make a function to turn all our data (x & y) into batches

```

# Define the batch size, 32 is a good start
BATCH_SIZE = 32

# Create a function to turn data into batches
def create_data_batches(X, y=None, batch_size=BATCH_SIZE, valid_data=False, test_data=False):
    """
    Creates batches of data out of image (X) and label (y) pairs.
    Shuffles the data if it's training data but doesn't shuffle if it's validation data.
    Also accepts test data as input (no labels).
    """
    # If the data is a test dataset, we probably don't have have labels
    if test_data:
        print("Creating test data batches...")
        data = tf.data.Dataset.from_tensor_slices((tf.constant(X))) # only filepaths (no label)
        data_batch = data.map(process_image).batch(BATCH_SIZE)
        return data_batch

```

```

# If the data is a valid dataset, we don't need to shuffle it
elif valid_data:
    print("Creating validation data batches...")
    data = tf.data.Dataset.from_tensor_slices((tf.constant(X), # filepaths
                                                tf.constant(y))) # labels
    data_batch = data.map(get_image_label).batch(BATCH_SIZE)
    return data_batch

else:
    print("Creating training data batches...")
    # Turn filepaths and labels into Tensors
    data = tf.data.Dataset.from_tensor_slices((tf.constant(X),
                                                tf.constant(y)))
    # Shuffling pathnames and labels before mapping image processor function is faster than
    data = data.shuffle(buffer_size=len(X))

    # Create (image, label) tuples (this also turns the iamge path into a preprocessed image)
    data = data.map(get_image_label)

    # Turn the training data into batches
    data_batch = data.batch(BATCH_SIZE)
    return data_batch

# Create training and validation data batches
train_data = create_data_batches(X_train, y_train)
val_data = create_data_batches(X_val, y_val, valid_data=True)

Creating training data batches...
Creating validation data batches...

# Check out the different attributes of our data batches
train_data.element_spec, val_data.element_spec

((TensorSpec(shape=(None, 224, 224, 3), dtype=tf.float32, name=None),
  TensorSpec(shape=(None, 120), dtype=tf.bool, name=None)),
 (TensorSpec(shape=(None, 224, 224, 3), dtype=tf.float32, name=None),
  TensorSpec(shape=(None, 120), dtype=tf.bool, name=None)))

```

▼ Visualizing Data batches

our data is now in batches, however these can be a little hard so lets visualize

```

import matplotlib.pyplot as plt

# Create a function for viewing images in a data batch
def show_25_images(images, labels):
    """
    Displays a plot of 25 images and their labels from a data batch.
    """

    # Setup the figure
    plt.figure(figsize=(10, 10))

```

```
# Loop through 25 (for displaying 25 images)
for i in range(25):
    # Create subplots (5 rows, 5 columns)
    ax = plt.subplot(5, 5, i+1)
    # Display an image
    plt.imshow(images[i])
    # Add the image label as the title
    plt.title(unique_breeds[labels[i].argmax()])
    # Turn the grid lines off
    plt.axis("off")

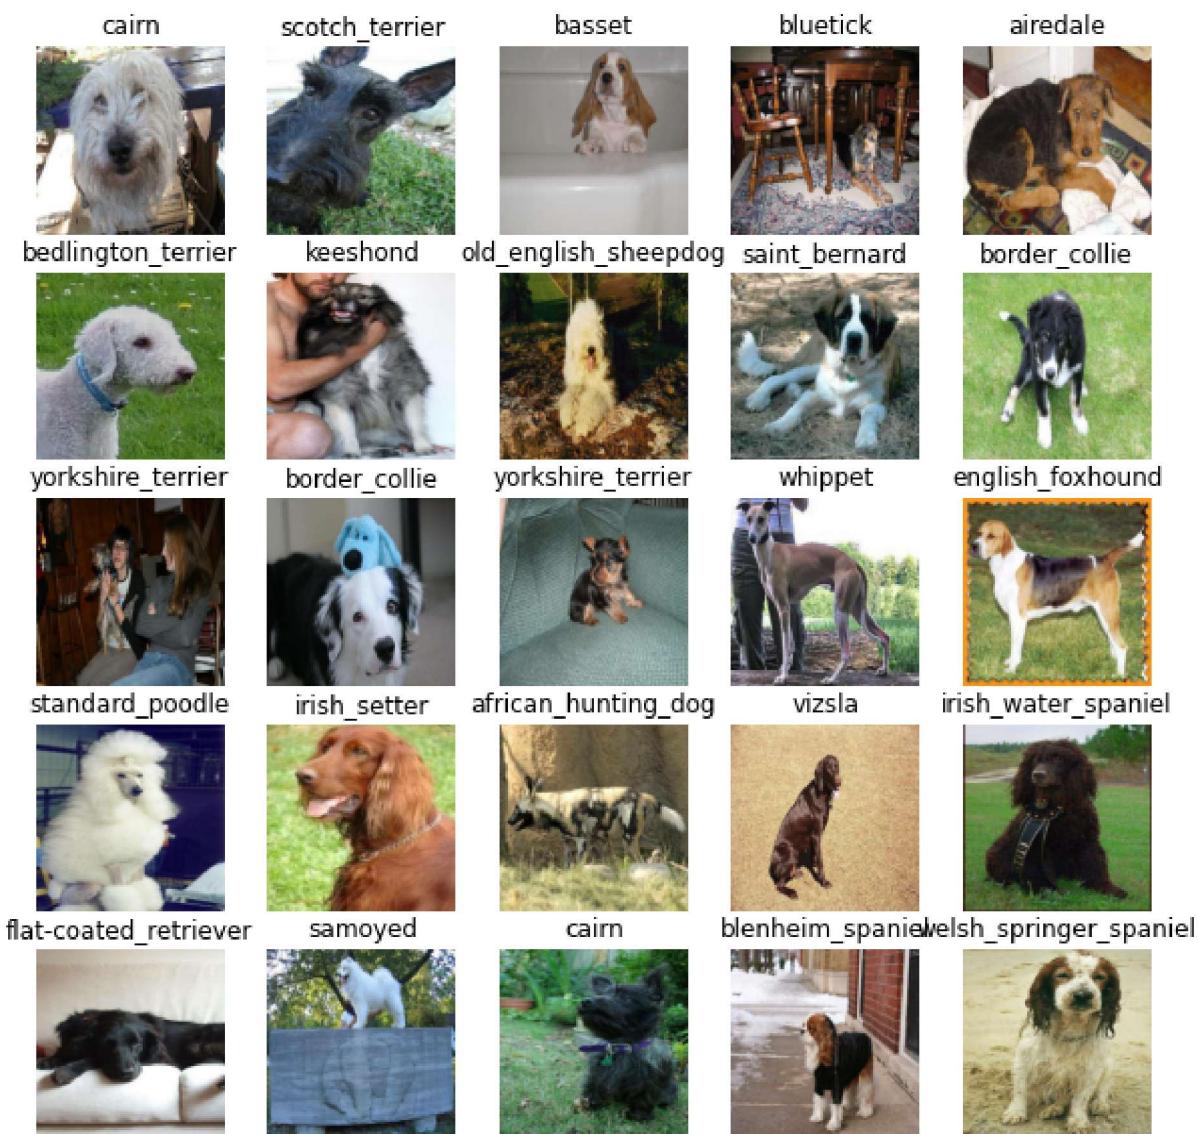
train_data

<BatchDataset element_spec=(TensorSpec(shape=(None, 224, 224, 3), dtype=tf.float32,
name=None), TensorSpec(shape=(None, 120), dtype=tf.bool, name=None))>

#train_images, train_labels = next(train_data.as_numpy_iterator())
#len(train_images), len(train_labels)

# now let's visualize the data in a training batch
train_images, train_labels = next(train_data.as_numpy_iterator())
show_25_images(train_images, train_labels)
```

```
#now let visualize our validation set
val_images, val_labels = next(val_data.as_numpy_iterator())
show_25_images(val_images, val_labels)
```



▼ Building a model

Before we build a model, there are a few things we need to define:

- The input shape (our images shape, in the form of Tensors) to our model.
- The output shape (image labels, in the form of Tensors) of our model.
- The URL of the model we want to use from TensorFlow Hub -

https://tfhub.dev/google/imagenet/mobilenet_v2_130_224/classification/4

IMG_SIZE

224

```

# Setup input shape to the model
INPUT_SHAPE = [None, IMG_SIZE, IMG_SIZE, 3] # batch, height, width, colour channels

# Setup output shape of our model
OUTPUT_SHAPE = len(unique_breeds)

# Setup model URL from TensorFlow Hub
MODEL_URL = "https://tfhub.dev/google/imagenet/mobilenet_v2_130_224/classification/5"

```

INPUT_SHAPE

[None, 224, 224, 3]

Now we've got our input, output and model ready, let's put them all together into Keras deep learning model.

let's create function which:

- Takes the input space and model we choose the parameter.
- Define layers in akerel model we've chosen as parameters.
- Define the layers in kerral (do this first, them this)
- Build the model
- Return the model

```

# Create a function which builds a Keras model
def create_model(input_shape=INPUT_SHAPE, output_shape=OUTPUT_SHAPE, model_url=MODEL_URL):
    print("Building model with:", MODEL_URL)

    # Setup the model layers
    model = tf.keras.Sequential([
        hub.KerasLayer(MODEL_URL), # Layer 1 (input layer)
        tf.keras.layers.Dense(units=output_shape,
                             activation="softmax") # Layer 2 (output layer)
    ])

    # Compile the model
    model.compile(
        loss=tf.keras.losses.CategoricalCrossentropy(),
        optimizer=tf.keras.optimizers.Adam(),
        metrics=["accuracy"]
    )

    # Build the model
    model.build(INPUT_SHAPE)

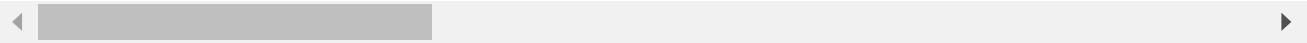
    return model

model = create_model()
model.summary()

```

```
WARNING:tensorflow:Please fix your imports. Module tensorflow.python.training.tracker
Building model with: https://tfhub.dev/google/imagenet/mobilenet\_v2\_130\_224/classification
WARNING:tensorflow:From /usr/local/lib/python3.8/dist-packages/tensorflow/python/autograph/contrib/keras.py:107: 
Instructions for updating:
Lambda functions will be no more assumed to be used in the statement where they are used.
Model: "sequential"
```

Layer (type)	Output Shape	Param #
<hr/>		
keras_layer (KerasLayer)	(None, 1001)	5432713
dense (Dense)	(None, 120)	120240
<hr/>		
Total params: 5,552,953		
Trainable params: 120,240		
Non-trainable params: 5,432,713		



▼ Creating callbacks

callbacks are helper function a model can use during training to do such things as save its progress, check its progress or stop training early if a model stop improving

we'll create callbacks one for TensorBoard which helps track our models progress and another for early stopping which prevents our model from training for too long

TensorBoard callback

To setup a tensorboard callback we need to set 3 things

1. Load the TensorBoard notebook extension
2. create a tensorBoard callback which is able to save logs to a directory and pass it our model fit() function
3. visualise our models training logs with % tensorboard magic function

```
# Load TensorBoard notebook extension
%load_ext tensorboard

import datetime

# Create a function to build a TensorBoard callback
def create_tensorboard_callback():
    # Create a log directory for storing TensorBoard logs
    logdir = os.path.join("drive/MyDrive/dog Vision/logs",
                          # Make it so the logs get tracked whenever we run an experiment
                          datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
    return tf.keras.callbacks.TensorBoard(logdir)
```

Early stopping callback early stopping helps our model from overfitting by stopping training if a certain evsuation metric stop improving

```
#create early stoping callback
early_stopping = tf.keras.callbacks.EarlyStopping(monitor="val_accuracy",
                                                    patience=3)
```

▼ Training a model(on subset of data)

our first model is only going to train on 1000 images to make sure everything is working

```
NUM_EPOCHS = 100 #@param {type:"slider", min:10, max:100, step:10}
```

100 

```
# check to make sure we're still running on a gpu
print("Gpu", "available (yes)" if tf.config.list_physical_devices("GPU") else "not available")

Gpu available (yes)
```

Let's create a funtion which train a model.

- Create a model using `create_model()`
- setup a TensorBoard callback using `create_tensorboard_callback()`
- call the `fit()` funtion on our model passing it the training data, validation data, number of epochs to train for(`NUM_EPOCHS`) and the callbacks we'd like to use
- return the model

```
# Build a function to train and return a trained model
def train_model():
    """
    Trains a given model and returns the trained version.
    """
    # Create a model
    model = create_model()

    # Create new TensorBoard session everytime we train a model
    tensorboard = create_tensorboard_callback()

    # Fit the model to the data passing it the callbacks we created
    model.fit(x=train_data,
               epochs=NUM_EPOCHS,
               validation_data=val_data,
               validation_freq=1, # check validation metrics every epoch
               callbacks=[tensorboard, early_stopping])

    return model
```

```
# Fit the model to the data
model = train_model()

Building model with: https://tfhub.dev/google/imagenet/mobilenet\_v2\_130\_224/classification
Epoch 1/100
25/25 [=====] - 118s 3s/step - loss: 4.6866 - accuracy: 0.0%
Epoch 2/100
25/25 [=====] - 7s 272ms/step - loss: 1.6997 - accuracy: 0.6%
Epoch 3/100
25/25 [=====] - 4s 152ms/step - loss: 0.5775 - accuracy: 0.9%
Epoch 4/100
25/25 [=====] - 7s 267ms/step - loss: 0.2574 - accuracy: 0.9%
Epoch 5/100
25/25 [=====] - 6s 228ms/step - loss: 0.1487 - accuracy: 0.9%
Epoch 6/100
25/25 [=====] - 5s 190ms/step - loss: 0.1001 - accuracy: 0.9%
Epoch 7/100
25/25 [=====] - 3s 124ms/step - loss: 0.0757 - accuracy: 1.0%
Epoch 8/100
25/25 [=====] - 6s 233ms/step - loss: 0.0606 - accuracy: 1.0%
Epoch 9/100
25/25 [=====] - 4s 168ms/step - loss: 0.0494 - accuracy: 1.0%
Epoch 10/100
25/25 [=====] - 6s 256ms/step - loss: 0.0416 - accuracy: 1.0%
Epoch 11/100
25/25 [=====] - 6s 231ms/step - loss: 0.0358 - accuracy: 1.0%
Epoch 12/100
25/25 [=====] - 4s 151ms/step - loss: 0.0313 - accuracy: 1.0%
Epoch 13/100
25/25 [=====] - 5s 200ms/step - loss: 0.0276 - accuracy: 1.0%
Epoch 14/100
25/25 [=====] - 5s 186ms/step - loss: 0.0246 - accuracy: 1.0%
Epoch 15/100
25/25 [=====] - 4s 171ms/step - loss: 0.0221 - accuracy: 1.0%
Epoch 16/100
25/25 [=====] - 5s 187ms/step - loss: 0.0200 - accuracy: 1.0%
Epoch 17/100
25/25 [=====] - 3s 122ms/step - loss: 0.0182 - accuracy: 1.0%
Epoch 18/100
25/25 [=====] - 3s 118ms/step - loss: 0.0167 - accuracy: 1.0%
Epoch 19/100
25/25 [=====] - 4s 171ms/step - loss: 0.0153 - accuracy: 1.0%
Epoch 20/100
25/25 [=====] - 3s 119ms/step - loss: 0.0142 - accuracy: 1.0%
Epoch 21/100
25/25 [=====] - 4s 150ms/step - loss: 0.0131 - accuracy: 1.0%
```



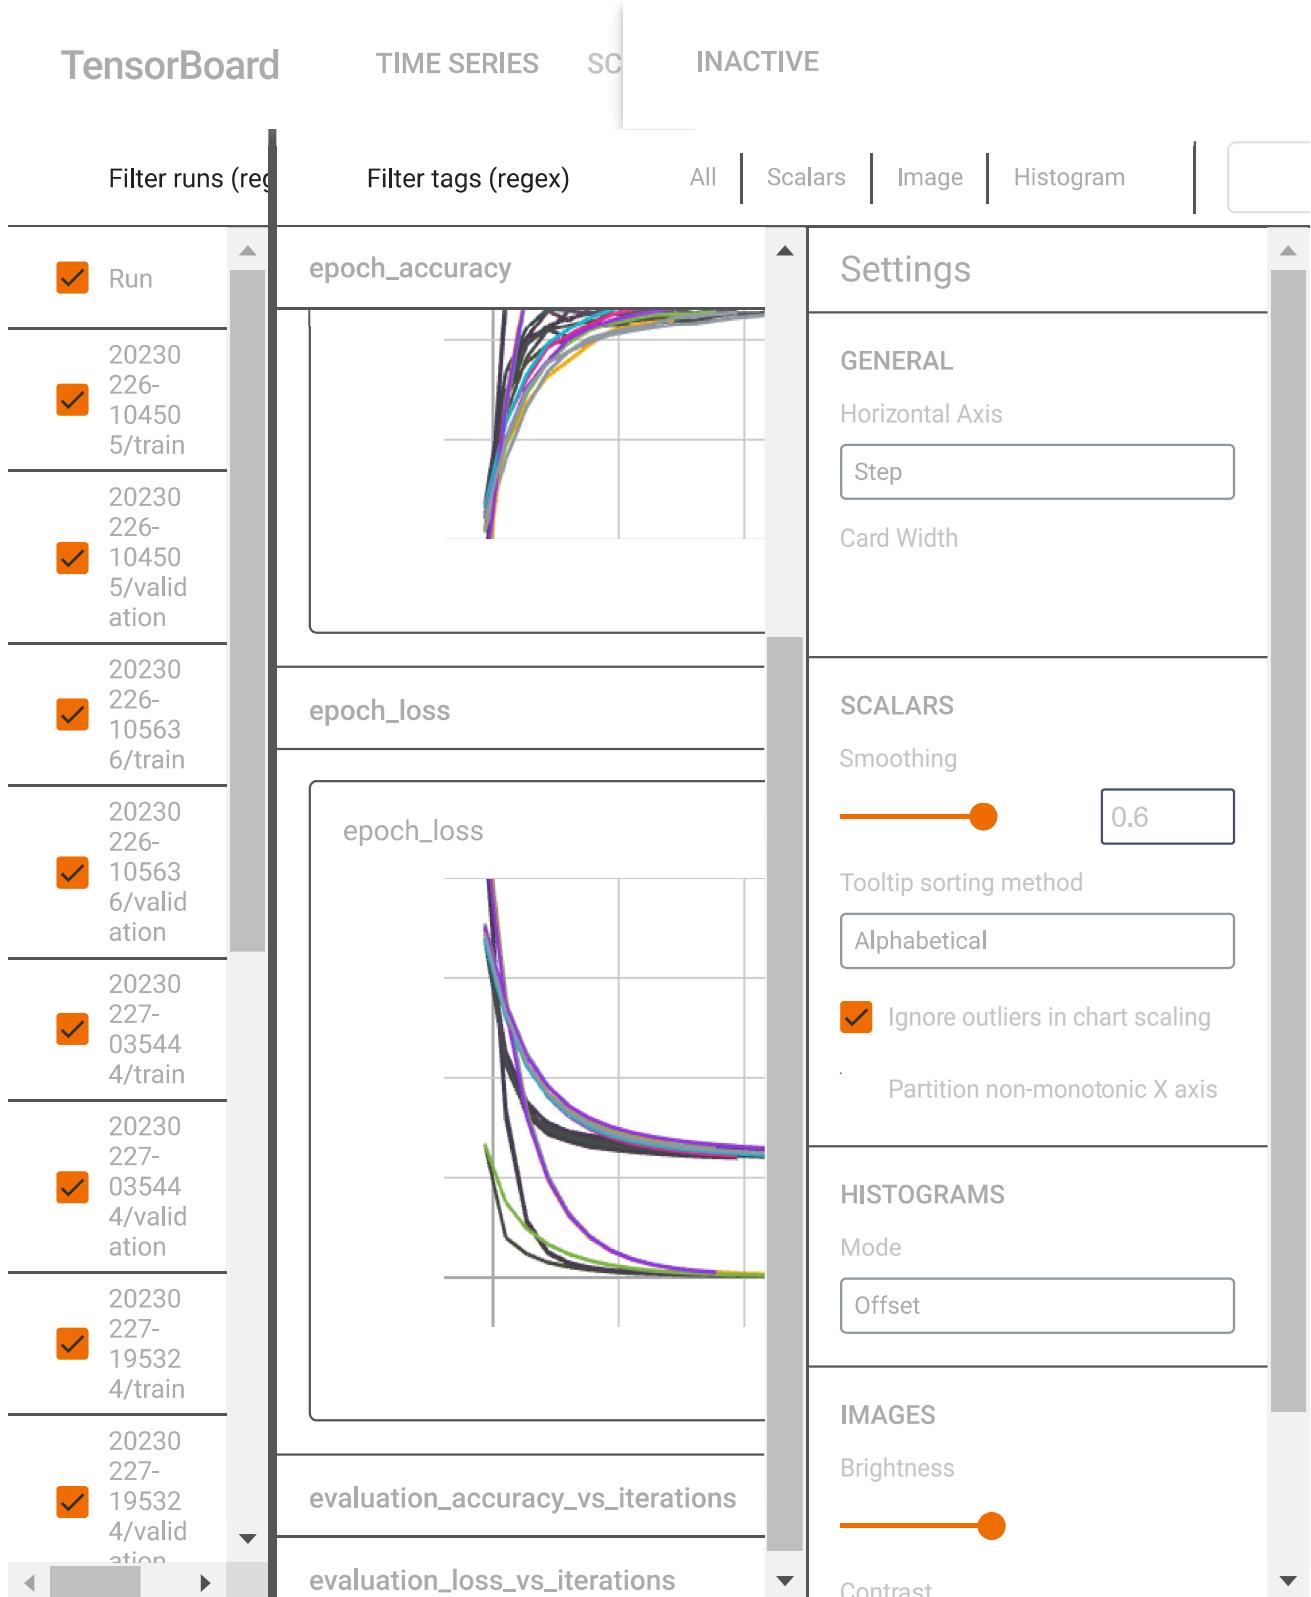
Question: it looks like our model is overfitting because it's performing better on the training dataset, now we have to prevent this

Note: overfitting is good things it mean our model is learning

▼ Checking the TensorBoard logs

The tensorboard magic function (%tensorboard) will access the logs directory we created earlier and visualize its

```
%tensorboard --logdir drive/MyDrive/dog\ Vision/logs
```



- ▼ Making and evaluation prediction using a trained mode

```
val_data

<BatchDataset element_spec=(TensorSpec(shape=(None, 224, 224, 3), dtype=tf.float32,
name=None), TensorSpec(shape=(None, 120), dtype=tf.bool, name=None))>

#Make prediction on the validation data( not used to train)
predictions = model.predict(val_data, verbose=1)
predictions

7/7 [=====] - 2s 110ms/step
array([[5.5887608e-04, 3.7422622e-05, 4.3527267e-04, ..., 9.6607349e-05,
       3.7021960e-05, 3.6021501e-03],
       [5.8099249e-04, 4.4168308e-04, 7.1741315e-03, ..., 5.0537736e-04,
       8.6551317e-04, 1.1108863e-04],
       [1.1914876e-05, 9.2963433e-05, 9.3570234e-05, ..., 2.3514256e-05,
       1.3531989e-05, 1.6326634e-05],
       ...,
       [4.0605960e-07, 3.6040146e-06, 1.3801128e-05, ..., 5.9686784e-07,
       1.0119476e-05, 9.9214958e-06],
       [2.6371100e-03, 1.3476177e-04, 3.0641491e-05, ..., 7.4622651e-05,
       1.7876535e-05, 1.5269596e-02],
       [2.8074608e-04, 5.0363928e-05, 2.2356566e-04, ..., 4.5520077e-03,
       3.6310891e-04, 2.4030908e-04]], dtype=float32)

predictions.shape

(200, 120)

np.sum(predictions[0])

0.99999994

np.sum(predictions[1])

1.0000001

len(y_val)

200

len(unique_breeds)

120

# First prediction
index=42
print(predictions[index])
print(f"Max value (probability of prediction): {np.max(predictions[index])}")
print(f"Sum: {np.sum(predictions[index])}")
print(f"Max index: {np.argmax(predictions[index])}")
print(f"prediction label: {unique_breeds[np.argmax(predictions[index])]}")
```

```
[8.20327987e-05 7.89824207e-06 4.29773354e-05 1.12899543e-05
 9.72048147e-04 2.53537019e-05 3.81863574e-05 9.00838582e-04
 1.72189600e-03 2.42623035e-02 3.57387689e-05 1.81031760e-06
 5.89359697e-05 1.37524924e-03 3.67526984e-04 1.87534699e-03
 6.19850289e-06 2.47546122e-04 2.84035486e-05 4.87918842e-05
 1.28817583e-05 2.11111968e-04 5.46215488e-05 5.67920324e-06
 6.00469345e-03 8.14471787e-06 1.17331356e-05 1.22872416e-05
 6.51717564e-05 1.65542970e-05 1.74517390e-05 1.49477455e-05
 6.58916679e-06 9.84527196e-06 1.24053495e-05 9.49125115e-06
 9.83546197e-05 6.92596150e-05 1.60686504e-05 1.40489146e-01
 5.06710312e-05 6.46327953e-06 9.12093790e-04 4.55272044e-07
 4.70150844e-04 8.26575779e-05 3.48368158e-05 4.50116495e-05
 7.26891221e-06 2.77275667e-05 2.91929555e-05 4.15743780e-05
 1.00616577e-04 3.77456145e-03 2.28306635e-05 1.41870201e-04
 9.40239261e-06 1.53879355e-05 1.66538339e-05 2.50135072e-05
 1.56810820e-05 3.48350062e-04 1.69014754e-06 1.11326035e-05
 2.91798169e-05 2.54306851e-05 1.25151710e-05 4.28605999e-04
 6.31184012e-05 3.87946011e-06 2.00269205e-05 1.51602071e-05
 1.76989270e-05 8.36382242e-05 7.92730716e-05 2.70709006e-05
 2.17459619e-05 4.99418275e-06 6.03963144e-06 1.55487956e-04
 1.40801853e-06 2.62454523e-05 2.70585624e-05 4.40572505e-04
 3.51065828e-04 4.61304489e-06 3.45284170e-05 1.10182623e-06
 5.60571016e-06 3.80217971e-04 1.06786953e-04 7.95481128e-07
 5.76259568e-04 5.34872815e-05 1.09417324e-05 1.83027696e-05
 1.61075568e-05 1.08493168e-05 6.91577316e-06 2.64678729e-05
 6.76083482e-06 1.47478413e-05 4.15795730e-05 3.30648472e-05
 1.11242683e-04 3.45200206e-05 7.03897531e-05 1.38518760e-06
 8.68570132e-05 1.10214358e-04 1.60305826e-05 3.11137701e-04
 8.43346497e-05 8.09729338e-01 1.06234998e-04 1.60526673e-04
 2.05086108e-05 3.32247691e-06 4.71748179e-04 4.46096492e-05]
Max value (probability of prediction): 0.8097293376922607
Sum: 0.9999998807907104
Max index: 113
prediction label: walker_hound
```

Having the above functionality is great but we want to be able to do it scale and it would be even better if we could see the image the prediction is being made on

Note: prediction are also known as confidence levels

```
# Turn prediction probabilities into their respective label (easier to understand)
def get_pred_label(prediction_probabilities):
    """
    Turns an array of prediction probabilities into a label.
    """
    return unique_breeds[np.argmax(prediction_probabilities)]
```

```
# Get a predicted label based on an array of prediction probabilities
pred_label = get_pred_label(predictions[81])
pred_label
```

```
'dingo'
```

```

val_data

<BatchDataset element_spec=(TensorSpec(shape=(None, 224, 224, 3), dtype=tf.float32,
name=None), TensorSpec(shape=(None, 120), dtype=tf.bool, name=None))>

```

Since our validation data is batch data set so we have to unbatch it to make on the prediction on the validation images and the compare those prediction to the validation labels(truth label)

```

#create a function to unbatch the dataset
# Create a function to unbatch a batch dataset
def unbatchify(data):
    """
    Takes a batched dataset of (image, label) Tensors and returns separate arrays
    of images and labels.
    """

    images = []
    labels = []
    # Loop through unbatched data
    for image, label in data.unbatch().as_numpy_iterator():
        images.append(image)
        labels.append(unique_breeds[np.argmax(label)])
    return images, labels

# Unbatchify the validation data
val_images, val_labels = unbatchify(val_data)
val_images[0], val_labels[0]

(array([[ [0.29599646, 0.43284872, 0.3056691 ],
         [0.26635826, 0.32996926, 0.22846507],
         [0.31428418, 0.27701408, 0.22934894],
         ...,
         [0.77614343, 0.82320225, 0.8101595 ],
         [0.81291157, 0.8285351 , 0.8406944 ],
         [0.8209297 , 0.8263737 , 0.8423668 ]],

      [[0.2344871 , 0.31603682, 0.19543913],
       [0.3414841 , 0.36560842, 0.27241898],
       [0.45016077, 0.40117094, 0.33964607],
       ...,
       [0.7663987 , 0.8134138 , 0.81350833],
       [0.7304248 , 0.75012016, 0.76590735],
       [0.74518913, 0.76002574, 0.7830809 ]],


      [[0.30157745, 0.3082587 , 0.21018331],
       [0.2905954 , 0.27066195, 0.18401104],
       [0.4138316 , 0.36170745, 0.2964005 ],
       ...,
       [0.79871625, 0.8418535 , 0.8606443 ],
       [0.7957738 , 0.82859945, 0.8605655 ],
       [0.75181633, 0.77904975, 0.8155256 ]],


      ...

```

```

[[0.9746779 , 0.9878955 , 0.9342279 ],
[0.99153054, 0.99772066, 0.9427856 ],
[0.98925114, 0.9792082 , 0.9137934 ],
...,
[0.0987601 , 0.0987601 , 0.0987601 ],
[0.05703771, 0.05703771, 0.05703771],
[0.03600177, 0.03600177, 0.03600177]],

[[0.98197854, 0.9820659 , 0.9379411 ],
[0.9811992 , 0.97015417, 0.9125648 ],
[0.9722316 , 0.93666023, 0.8697186 ],
...,
[0.09682598, 0.09682598, 0.09682598],
[0.07196062, 0.07196062, 0.07196062],
[0.0361607 , 0.0361607 , 0.0361607 ]],

[[0.97279435, 0.9545954 , 0.92389745],
[0.963602 , 0.93199134, 0.88407487],
[0.9627158 , 0.91253304, 0.8460338 ],
...,
[0.08394483, 0.08394483, 0.08394483],
[0.0886985 , 0.0886985 , 0.0886985 ],
[0.04514172, 0.04514172, 0.04514172]]], dtype=float32), 'cairn')

get_pred_label(val_labels[0])
'affenpinscher'

```

Now we've got ways to get

- prediction label
- validation label
- validation images Let make some funtio to visualise it

we crate funtion which:

- take an array prediction probablites, an array of truth labels and an array of images and iterator
- convert the prediction probablites to a predicted label
- plot the prediction probablites, the truth label and the target image on a single plot

```

def plot_pred(prediction_probabilities, labels, images, n=1):
    """
    View the prediction, ground truth and image for sample n
    """
    pred_prob, true_label, image = prediction_probabilities[n], labels[n], images[n]

    # Get the pred label
    pred_label = get_pred_label(pred_prob)

    # Plot image & remove ticks
    plt.imshow(image)
    plt.xticks([])

```

```

plt.yticks([])

# Change the colour of the title depending on if the prediction is right or wrong
if pred_label == true_label:
    color = "green"
else:
    color = "red"

# Change plot title to be predicted, probability of prediction and truth label
plt.title("{} {:.2f}% {}".format(pred_label,
                                  np.max(pred_prob)*100,
                                  true_label),
          color=color)

plot_pred(prediction_probabilities=predictions,
           labels=val_labels,
           images=val_images,
           n=10)

```

`yorkshire_terrier 35% yorkshire_terrier`



Now we've got one function to visualize our models top prediction, let's make another to view our models top 10 predictions.

This function will:

- Take an input of prediction probabilities array and a ground truth array and an integer
- Find the prediction using `get_pred_label()`
- Find the top 10: Prediction probabilities indexes

Prediction probabilities values

Prediction labels

- Plot the top 10 prediction probability values and labels, coloring the true label green

```

def plot_pred_conf(prediction_probabilities, labels, n=1):
    """
    plus the top 10 highest preddiction confidence along with the truth label for sample n
    """

```

```

pred_prob, true_label = prediction_probabilities[n], labels[n]

#Get the prediction label
pred_label = get_pred_label(pred_prob)

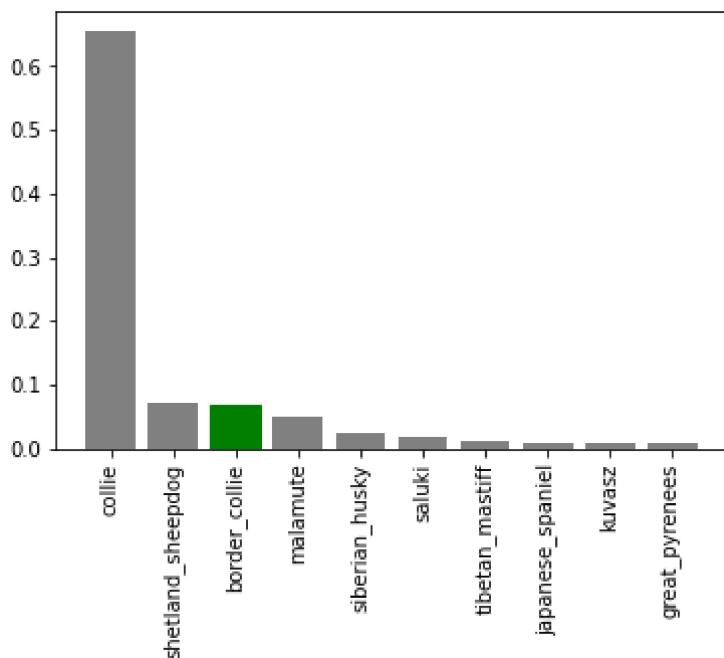
# find the top 10 prediction confidence indexes
top_10_pred_indexes = pred_prob.argsort()[-10:][::-1]
# find the top 10 prediction confidence values
top_10_pred_values = pred_prob[top_10_pred_indexes]
# find the top 10 prediction labels
top_10_pred_labels = unique_breeds[top_10_pred_indexes]

# set up the plot
top_plot = plt.bar(np.arange(len(top_10_pred_labels)),
                   top_10_pred_values,
                   color="grey")
plt.xticks(np.arange(len(top_10_pred_labels)),
           labels=top_10_pred_labels,
           rotation="vertical")

#change color of true label
if np.isin(true_label, top_10_pred_labels):
    top_plot[np.argmax(top_10_pred_labels == true_label)].set_color("green")
else:
    pass

plot_pred_conf(prediction_probabilities=predictions,
                labels=val_labels,
                n=9)

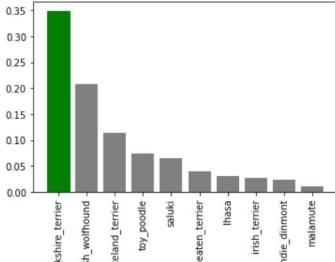
```



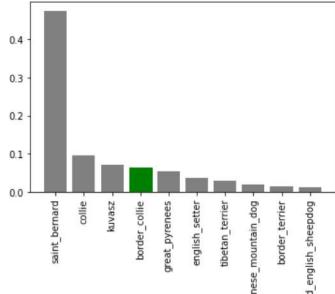
Now we've got some function to help us visualize our prediction and evaluate our model let's check out few

```
# Let's check out a few predictions and their different values
i_multiplier = 10
num_rows = 3
num_cols = 2
num_images = num_rows*num_cols
plt.figure(figsize=(10*num_cols, 5*num_rows))
for i in range(num_images):
    plt.subplot(num_rows, 2*num_cols, 2*i+1)
    plot_pred(prediction_probabilities=predictions,
               labels=val_labels,
               images=val_images,
               n=i+i_multiplier)
    plt.subplot(num_rows, 2*num_cols, 2*i+2)
    plot_pred_conf(prediction_probabilities=predictions,
                   labels=val_labels,
                   n=i+i_multiplier)
plt.tight_layout(h_pad=1.0)
plt.show()
```

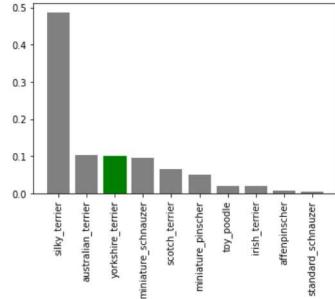
yorkshire_terrier 35% yorkshire_terrier



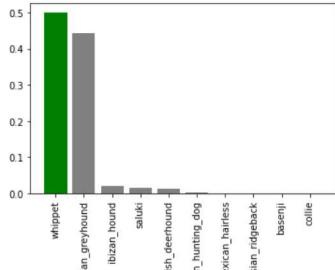
saint_bernard 47% border_collar



silky_terrier 49% yorkshire_terrier



whippet 50% whippet



english_foxhound 77% english_foxhound



miniature_poodle 48% standard_poodle



▼ Saving and reloading a trained model

```
# create a function to save a model
def save_model(model, suffix=None):
    """
    Save a given model in a models directory and appends a sufix (string).
    """

    # create a model directory pathname with current time
    modeldir = os.path.join("drive/MyDrive/dog Vision/models",
                           datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
    model_path = modeldir + " " + suffix + ".h5" # save formate of model
    print(f"Saving model to: {model_path}...")
    model.save(model_path)
    return model_path

# Create a function to load a trained model
def load_model(model_path):
    """
    Loads a saved model from a specified path.
    """

    print(f"Loading saved model from: {model_path}")
    model = tf.keras.models.load_model(model_path,
                                       custom_objects={"KerasLayer":hub.KerasLayer})
    return model

#save our model trained on 1000 images
save_model(model, suffix="1000-images-mobilenetv2-Adam")
```

```
Saving model to: drive/MyDrive/dog Vision/models/20230301-190651_1000-images-mobilenet_v2.h5
# load a trained model
loaded_1000_image_model = load_model('drive/MyDrive/dog Vision/models/20230228-205458_1000-images-mobilenet_v2.h5')

Loading saved model from: drive/MyDrive/dog Vision/models/20230228-205458_1000-images-mobilenet_v2.h5
```

◀ ▶

```
# Evaluate the pre-saved model
model.evaluate(val_data)
```

```
7/7 [=====] - 1s 83ms/step - loss: 1.2047 - accuracy: 0.6756
[1.2046617269515991, 0.675000011920929]
```

◀ ▶

```
#evaluate the loaded model
loaded_1000_image_model.evaluate(val_data)
```

```
7/7 [=====] - 2s 130ms/step - loss: 1.2663 - accuracy: 0.6856
[1.2662937641143799, 0.6850000023841858]
```

▼ Trainig a big dog model on full data

```
len(X), len(y)
```

```
(10222, 10222)
```

```
X[:10]
```

```
['drive/MyDrive/dog Vision/train/000bec180eb18c7604dcecc8fe0dba07.jpg',
 'drive/MyDrive/dog Vision/train/001513dfcb2ffafc82cccf4d8bbaba97.jpg',
 'drive/MyDrive/dog Vision/train/001cdf01b096e06d78e9e5112d419397.jpg',
 'drive/MyDrive/dog Vision/train/00214f311d5d2247d5dfe4fe24b2303d.jpg',
 'drive/MyDrive/dog Vision/train/0021f9ceb3235effd7fcde7f7538ed62.jpg',
 'drive/MyDrive/dog Vision/train/002211c81b498ef88e1b40b9abf84e1d.jpg',
 'drive/MyDrive/dog Vision/train/00290d3e1fdd27226ba27a8ce248ce85.jpg',
 'drive/MyDrive/dog Vision/train/002a283a315af96eaea0e28e7163b21b.jpg',
 'drive/MyDrive/dog Vision/train/003df8b8a8b05244b1d920bb6cf451f9.jpg',
 'drive/MyDrive/dog Vision/train/0042188c895a2f14ef64a918ed9c7b64.jpg']
```

```
len(X_train)
```

```
800
```

```
# create a data batch with the full data set
full_data = create_data_batches(X, y)
```

```
Creating training data batches...
```

```
full_data

<BatchDataset element_spec=(TensorSpec(shape=(None, 224, 224, 3), dtype=tf.float32,
name=None), TensorSpec(shape=(None, 120), dtype=tf.bool, name=None))>

full_model = create_model()

Building model with: https://tfhub.dev/google/imagenet/mobilenet\_v2\_130\_224/classification
```



```
#create full model callbacks
full_model_tensorboard = create_tensorboard_callback()
# No validation set when training on all the data so we can't monitor validation accuracy
full_model_early_stopping = tf.keras.callbacks.EarlyStopping(monitor="accuracy",
                                                               patience=3)
```

this cell may be more time to complete

```
# Fit the full model to the full data
full_model.fit(x=full_data,
                epochs=NUM_EPOCHS,
                callbacks=[full_model_tensorboard, full_model_early_stopping])

Epoch 1/100
320/320 [=====] - 44s 124ms/step - loss: 1.3366 - accuracy:
Epoch 2/100
320/320 [=====] - 33s 103ms/step - loss: 0.3972 - accuracy:
Epoch 3/100
320/320 [=====] - 52s 160ms/step - loss: 0.2387 - accuracy:
Epoch 4/100
320/320 [=====] - 34s 105ms/step - loss: 0.1558 - accuracy:
Epoch 5/100
320/320 [=====] - 45s 140ms/step - loss: 0.1073 - accuracy:
Epoch 6/100
320/320 [=====] - 44s 137ms/step - loss: 0.0770 - accuracy:
Epoch 7/100
320/320 [=====] - 38s 119ms/step - loss: 0.0591 - accuracy:
Epoch 8/100
320/320 [=====] - 39s 121ms/step - loss: 0.0462 - accuracy:
Epoch 9/100
320/320 [=====] - 37s 114ms/step - loss: 0.0381 - accuracy:
Epoch 10/100
320/320 [=====] - 44s 139ms/step - loss: 0.0309 - accuracy:
Epoch 11/100
320/320 [=====] - 43s 135ms/step - loss: 0.0264 - accuracy:
Epoch 12/100
320/320 [=====] - 40s 124ms/step - loss: 0.0219 - accuracy:
Epoch 13/100
320/320 [=====] - 36s 112ms/step - loss: 0.0211 - accuracy:
Epoch 14/100
320/320 [=====] - 33s 104ms/step - loss: 0.0177 - accuracy:
Epoch 15/100
320/320 [=====] - 34s 105ms/step - loss: 0.0183 - accuracy:
<keras.callbacks.History at 0x7f4af06a75b0>
```

```
save_model(full_model, suffix ="full-image-set-mobilenetv2-Adam")  
Saving model to: drive/MyDrive/dog Vision/models/20230301-191659_full-image-set-mobil  
e/MyDrive/dog Vision/models/20230301-191659_full-image-set-mobilenetv2-Adam.h5'  
  
#load the full model  
loaded_full_model = load_model('drive/MyDrive/dog Vision/models/20230228-212927_full-image  
Loading saved model from: drive/MyDrive/dog Vision/models/20230228-212927_full-image-
```

▼ Making prediction on test dataset

since our model has been trained on images in the form of Tensor batches to make predictions on the test data, we'll have to get it into the same formate as we created `data_batches()` earlier which can take a list of filename as input and convert into tensor batches To make predictions on the test data we'll:

- Get the image filenamess
- convert filenames into test data batches using `create_databatches()` and setting `test_data` parameter True as test data don't have labels
- Make prediction array by passing the test batches to the `predict()` method called on our model

```
# load the image filenames  
test_path ="drive/MyDrive/dog Vision/test/"  
test_filenames = [test_path + fname for fname in os.listdir(test_path)]  
test_filenames[:10]  
  
['drive/MyDrive/dog Vision/test/e427b9e1ab1b7f09cfb02ac073f56f2d.jpg',  
 'drive/MyDrive/dog Vision/test/e421e13d0d4c977ab89fa9965221809c.jpg',  
 'drive/MyDrive/dog Vision/test/e30a97eb637239e3b15c903529e04800.jpg',  
 'drive/MyDrive/dog Vision/test/ddc0134def0da0472e55271da8b57c44.jpg',  
 'drive/MyDrive/dog Vision/test/e411a1b3681604f6321af7cb8e8f2de7.jpg',  
 'drive/MyDrive/dog Vision/test/e334f758c7944df19c98d49498d28c64.jpg',  
 'drive/MyDrive/dog Vision/test/deaf748bed865ec76758f086db04bfac.jpg',  
 'drive/MyDrive/dog Vision/test/e103930a8a4416f076c81a51907a7512.jpg',  
 'drive/MyDrive/dog Vision/test/e278b2f68fa5d146e0bf06eee5faa0e3.jpg',  
 'drive/MyDrive/dog Vision/test/e71d9f32ea6eb1c2d944e8f4e811c209.jpg']
```

```
len(test_filenames)
```

```
10357
```

```
#create test data batch  
test_data = create_data_batches(test_filenames, test_data=True)
```

```

Creating test data batches...

test_data

<BatchDataset element_spec=TensorSpec(shape=(None, 224, 224, 3), dtype=tf.float32,
name=None)>

# Make prediction on test data batch using the loaded full model
test_predictions = loaded_full_model.predict(test_data,
                                             verbose=1)

324/324 [=====] - 148s 445ms/step

# save prediction(numpy array) to csv file for accesss later
#np.savetxt("drive/MyDrive/dog Vision/preds_array.csv", test_predictions, delimiter=",")

# load predictions(numpy array) from csv file
test_predictions=np.loadtxt("drive/MyDrive/dog Vision/preds_array.csv", delimiter=",")

test_predictions[:10]

array([[5.04669773e-09, 8.73544170e-09, 6.07846107e-10, ...,
       7.51192372e-11, 9.99876499e-01, 4.93539787e-09],
      [2.35612148e-13, 6.02668249e-10, 1.78500215e-07, ...,
       7.32975947e-09, 1.21073871e-08, 1.12957907e-12],
      [1.84990611e-04, 1.08305764e-09, 1.03033591e-08, ...,
       6.73671394e-08, 1.00510517e-10, 1.62898202e-03],
      ...,
      [1.66965686e-09, 1.08444431e-06, 1.84264370e-09, ...,
       4.07181346e-08, 8.86748530e-05, 1.66152798e-08],
      [1.97171513e-02, 8.37505067e-05, 1.91008009e-08, ...,
       1.03585848e-10, 1.44219886e-10, 8.52987289e-01],
      [7.16603554e-10, 1.41188269e-04, 4.35292313e-05, ...,
       2.85418741e-06, 5.09059755e-04, 5.79477230e-04]]))

test_predictions.shape

(10357, 120)

```

▼ Making predictions on custom images we'll

- get the filepath of our own images
- Turn the filepath into data batches using `create_data_batches()` and since our images won't have labels we set the `test_data` parameter to true
- pass the custom image data batch to our model's `predict()` method
- convert the prediction output probabilities to prediction labels
- compare the predicted labels to the custom images

```

#get custom image filepaths
custom_path = "drive/MyDrive/dog Vision/my-dog-photo/"
custom_image_paths = [custom_path + fname for fname in os.listdir(custom_path)]


custom_image_paths

['drive/MyDrive/dog Vision/my-dog-photo/desi1.jpeg',
 'drive/MyDrive/dog Vision/my-dog-photo/german.jpeg']


#turn custom image into batch datasets
custom_data= create_data_batches(custom_image_paths, test_data=True)
custom_data

Creating test data batches...
<BatchDataset element_spec=TensorSpec(shape=(None, 224, 224, 3), dtype=tf.float32,
name=None)>


# make prediction on the custom data
custom_preds = loaded_full_model.predict(custom_data)

1/1 [=====] - 0s 39ms/step


custom_preds.shape

(2, 120)


# Get custom image prediction labels
custom_pred_labels = [get_pred_label(custom_preds[i]) for i in range (len(custom_preds))]
custom_pred_labels

['dingo', 'tibetan_mastiff']


# get custom images (our unbatchify() ) funtion won't work since aren't label we have to fi
custom_images=[]
# loop throug unbatched data
for image in custom_data.unbatch().as_numpy_iterator():
    custom_images.append(image)


import enum
# check custom image prediction
plt.figure (figsize=(10, 10))
for i, image in enumerate(custom_images):
    plt.subplot(1, 3, i+1)
    plt.xticks([])
    plt.yticks([])
    plt.title(custom_pred_labels[i])
    plt.imshow(image)

```

```
-----  
NameError Traceback (most recent call last)  
<ipython-input-1-5126b9ad9d64> in <module>  
      1 import enum  
      2 # check custom image prediction  
----> 3 plt.figure (figsize=(10, 10))  
      4 for i, image in enumerate(custom_images):  
      5     plt.subplot(1, 3, i+1)  
  
NameError: name 'plt' is not defined
```

